

Analysis

Producers & Consumers

Implementation:

Initial Construction:

Without skeleton code we had to decide how to construct the project in such a way that wouldn't be overly complicated without reason. We initially wanted to have four files; a main, a producer file, a consumer file and a product file. Ultimately, we decided on combining the producers, consumers and the main file for the sake of not having to pass too many arguments. This was due to POSIX threads requiring a single parameter as a void pointer. So too keep our solution simple we opted to use a few global variables.

Constructing Producers & First-Come-First-Serve:

In order to have the option to run consumers with Round Robin or First-Come-First-Serve scheduling we created two consumer functions. The structure of First-Come-First-Serve ended up very similar to that of our Producer function. Each Producer was given two binary mutexes, one for preventing other producers from changing the global variables and another specifically for our queue.

Round Robin:

For deciding how to construct our Round Robin consumer function we used a metaphorical real-world model of the problem. We said the Round Robin algorithm is like a group of people sharing snickers bars. The people being the producers, the snickers being the products but most importantly the bites they take represent the quanta.

Experimentation:

Zach wrote a python script using numpy and pyplot to test and graph the performance of our implementation.

How it works:

1. Runs two nested loops, to rotate through Quanta, then to test 5 times for each Quantum
2. Analyzes the stdout of the C++ code using a flag we built into the code to initiate test mode
3. Compares timing values and graphs the analyzed information
4. TO RUN THE PYTHON3 SCRIPT YOU MUST INSTALL NUMPY AND MATPLOTLIB

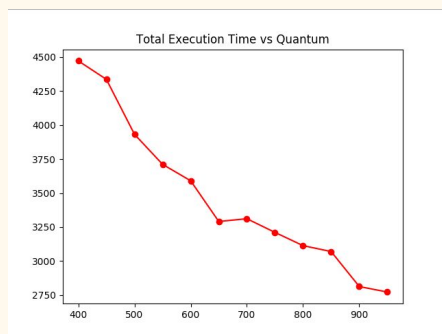
Expectation:

We expected to find a parabolic curve in the performance of the Round Robin meaning that there would be a “sweet spot” in quantum size. We also expected that quanta outside of this “sweet spot” would yield execution and turnaround times closer to, if not worse than, First-Come-First-Serve.

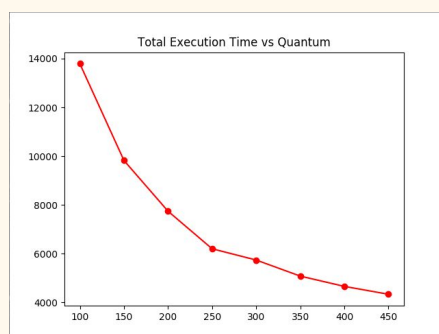
Results:

To our surprise increasing quanta only makes the Round Robin Scheduling more efficient. We assume this to be the case specifically with our Producers & Consumers problem, being that managing a single queue of objects is much simpler than full CPU scheduling. We hypothesize the main reason for the result varying from our initial expectation to be the fact that the threads we are working with have no I/O wait time embedded in them. In a real world scenario, threads would need to be interacting with the disk or waiting for user input which would cause heavy downfalls in the wait times, turnover times, and total execution time. Below are the graphs we produced to see these results. All graphs are as avg time of Y as Quantum increases.

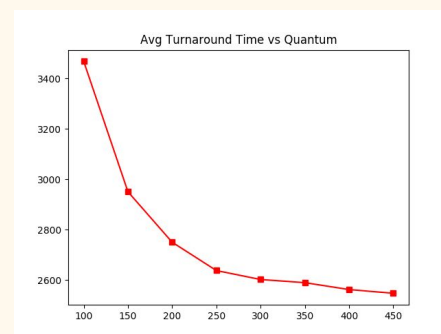
Input	Producer Throughput	RR Throughput	FCFS ThroughPut
10 10 10 10 x 10 10	10.161	540.85	10.1744
10 10 10 10 x 5000 10	10.2095	10.2105	Unchanged ^
10 10 100 10 0 10 10	439.711	4857.77	100.41
10 10 100 10 0 5000 10	10.0377	100.451	unchanged^



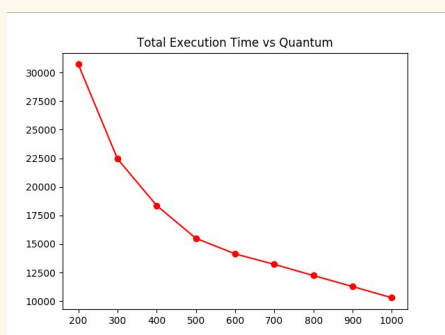
A



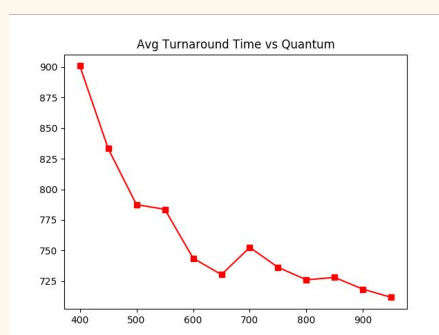
B



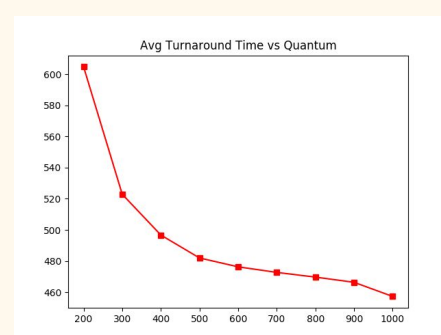
C



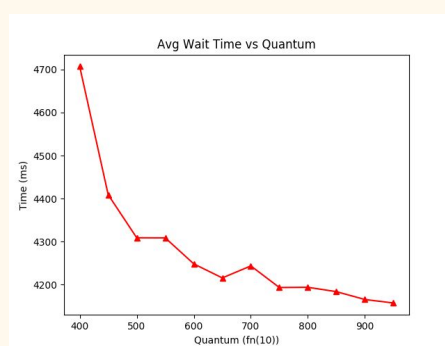
D



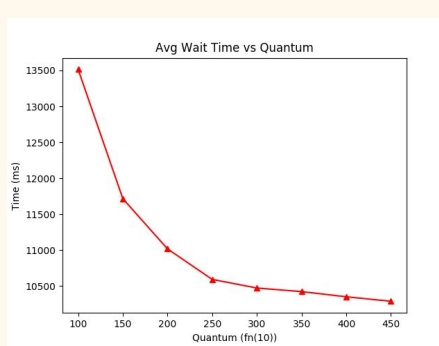
E



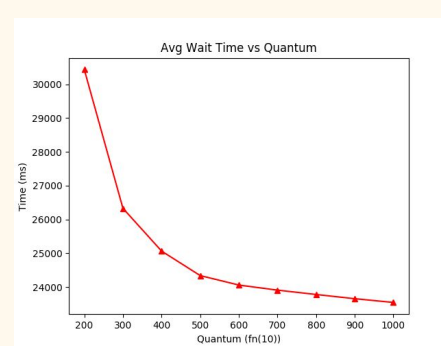
F



G



H



I